

# Self Supervised Learning



# Theme

---

SSL can uncover important information present in training data, and encode similarities/ dissimilarities

Can we learn building blocks of causal graphical model by SSL?

How much supervision do we need?

# Recent Work



---

- SSL with Data Augmentation (DA) provably isolates content from style [1]
  - Shows that SSL with DA learns a partition of the latent space into content and style variables
  - Style: latent codes which change when a data-point is transformed (by rotation, change in color saturation, camera angle, resizing,..)
  - Content remains invariant
  - Resembles a SCM where a Counterfactual is  $DA(x)$
  - Shows that we can uncover the causal structure from DA (to some degree)

# Content vs Style

## 1. How to differentiate Content vs Style variables?

- It is implicitly expressed from the data augmentation methods we apply
- The partition can be learned from the information contained in dataset
- The method assumes that content variables are invariant in a DA
- Employs a ContrastiveLearning framework (InfoNCE objective) to leverage the similarity between original and augmented datapoint



[1] Figure 2: (Left) Causal graph for the Causal3DIdent dataset. (Right) Two samples from each object class.

# Relation to Causal Factors

- Does identifying content variables in latent space always corresponds to causal factors?
  - The method shows good results in identifying content variables (invariant ones) from InfoNCE objective. By [1] Defn 4.1:

**Definition 4.1** (Block-identifiability). We say that the true content partition  $\mathbf{c} = \mathbf{f}^{-1}(\mathbf{x})_{1:n_c}$  is *block-identified* by a function  $\mathbf{g} : \mathcal{X} \rightarrow \mathcal{Z}$  if the inferred content partition  $\hat{\mathbf{c}} = \mathbf{g}(\mathbf{x})_{1:n_c}$  contains *all* and *only* information about  $\mathbf{c}$ , i.e., if there exists an *invertible* function  $\mathbf{h} : \mathbb{R}^{n_c} \rightarrow \mathbb{R}^{n_c}$  s.t.  $\hat{\mathbf{c}} = \mathbf{h}(\mathbf{c})$ .

- Used kernel ridge regression to predict ground truth  $\mathbf{c}$  and  $\mathbf{s}$  from the learnt representations  $\hat{\mathbf{c}} = \mathbf{g}(\mathbf{x})$  and report  $R^2$  coefficient. Highly accurate results on 3dldnet dataset when there is a causal dependency between style and content variables

Generative process			$R^2$ (nonlinear)	
p(chg.)	Stat.	Cau.	Content $\mathbf{c}$	Style $\mathbf{s}$
1.0	✗	✗	$1.00 \pm 0.00$	$0.07 \pm 0.00$
0.75	✗	✗	$1.00 \pm 0.00$	$0.06 \pm 0.05$
0.75	✓	✗	$0.99 \pm 0.00$	$0.02 \pm 0.04$
0.75	✓	✓	$0.96 \pm 0.00$	$0.00 \pm 0.00$

# Relation to Causal Factors

---

- Consider the image on right:
  - There is no recognizable object – a simple out-of-distribution sample
- Several other OOD examples: [link](#)
- **What if we attempt to learn the Content vs Style partition on a dataset of such images with their augmentations?**
  - Difficult to think what would be content or style
  - Perhaps no relation to learning ground truth latent causal factors
  - What information would an augmented pair help in encoding?



# Relation to Causal Factors

---

- A digression:
  - *This image is made by morphing several subjects such that the boundaries merge and diffuse erratically. This is an "ood" feature for human world where boundaries follow a distinctive rule of being continuous.*
  - *We can learn these very simple rules and associations from the given data itself, by exploiting large number of similarities and dissimilarities, between the given data*






# More with SSL

- My thoughts:
  - Contrastive Learning (CL) framework with data augmentation helps to uncover causal structure because common factors remain the same. The style attributes have "Content" variables as their Parents in SCM.
  - Along with DA, we can also utilize CL on patches of image.
    - What information can we get by pairing the two patches shown? For examples,
      - *With high probability, in this dataset,*

*similarities* → *an image can have sharp lines distinguishing two regions. A region can have a uniform solid color, or a tinge of multiple colors. (patch4 and patch3)*

*dissimilarities (counterfactuals)* → *Boundary lines can be sharp (patch1) or diffused (patch2)*





# More with SSL

---

- Facts and counterfactuals can help to learn distinguishing features
- Learning positive correlations between convolutional patches within a sample image can bring more information
  - Consider the example on right (Fig1): the identity of these icons is *a mic/camera with a slash on top*, although the underlying white icon is not connected.
  - Pairing the conv patches 1 and 2 will give a high similarity score, indicating that they belong to same context.





Fig 4.1: The icons for mic/camera off on Google Meet. Notice how we perceive it as a mic with a white slash on top, rather than two white scribbles separated by a red slant line

# More with SSL

---

Likewise, we can learn the probability of continuation of line segment from such patches within the dataset. This falls in line with Gestalt Principle of continuation.



**Figure 4.** Good continuation would suggest that we are more likely to perceive this as two overlapping lines, rather than four lines meeting in the center.



**Figure 5.** Closure suggests that we will perceive a complete circle and rectangle rather than a series of segments.

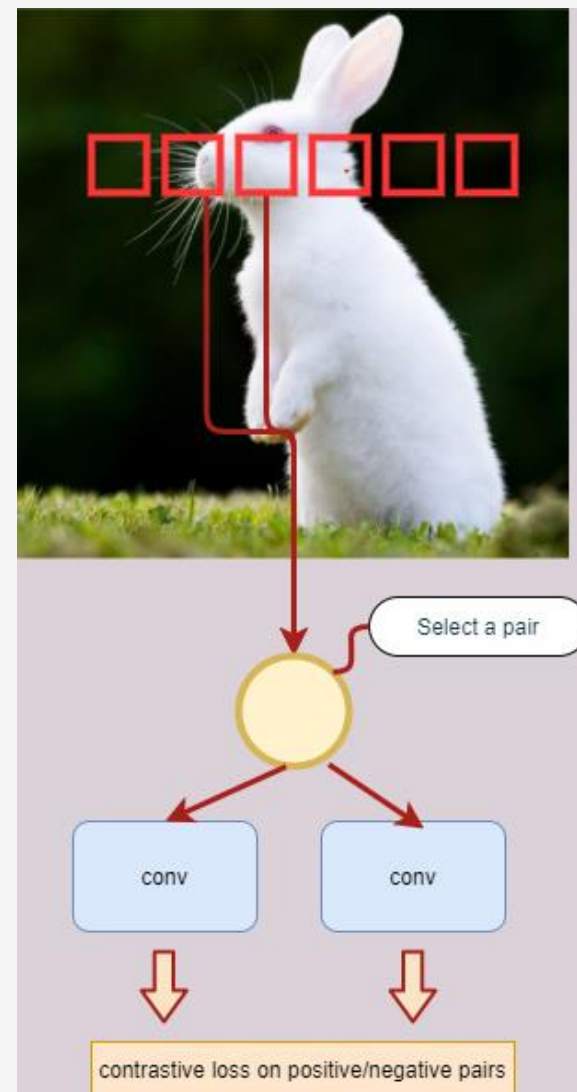
# How would this be useful?

---

- The way we can learn the probability of line segments being connected than disconnected, we can extract some information about various patches *without any supervision*.
- Key points:
  - Exploit similarities and dissimilarities between conv patch in dataset (images)
  - Repeat the above for different size of conv patches
  - Use already learned information to build a hierarchical structure?


# How would this be useful?

- Start with small window size
- Use a convNet to learn representations, followed by a CL loss
- How to decide a pair as positive/negative:
  - TBD
  - A simple threshold of distance between patch embeddings



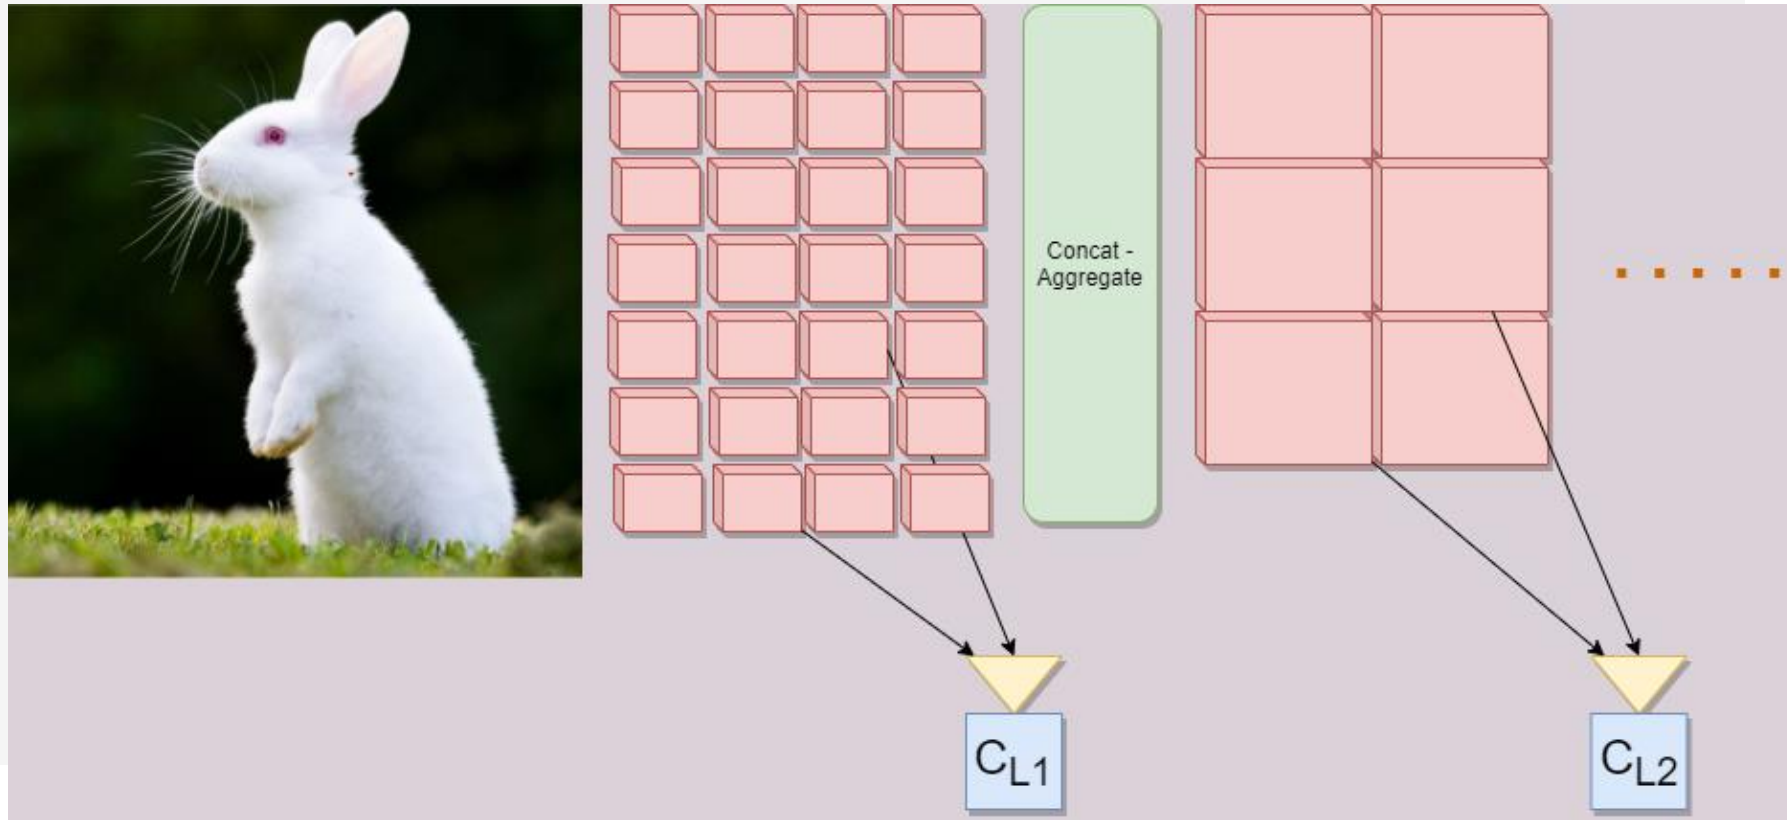
# How would this be useful?

- Iterate with increasing window size
- The representation learnt in previous step should be used to perform this iteration
- Stack these steps?



# How would this be useful?

- Joint learning:
  - At each step, create patches, choose a pair, pass through convNet, train on ContrastiveLoss ( $C_{Li}$ ) objective
  - At step\_K, concatenate embedding of constituent patches from step\_K-1 to use information learned in the previous step
  - Minimize  $Sum(C_{l1} + C_{l2} + \dots)$



# How would this be useful?

- Incremental learning:
  - At step\_K, train only on  $C_{Lk}$  objective, and freeze conv layers till step\_K-1. Train only the additional MLP while calculating  $C_L$

